

Vedant Desai

(248) 704 – 4852 | vedantde@umich.edu | github.com/Verdent06 | linkedin.com/in/vedantde06

EDUCATION

University of Michigan

B.S. in Computer Science and Economics

Expected May 2028

Ann Arbor, MI

- **GPA:** 3.66 / 4.0
- **Coursework:** Data Structures & Algorithms, Intro to Statistics and Data Analysis, Discrete Mathematics, Calculus III

EXPERIENCE

Michigan Data Consulting (MDC)

Jan 2026 – May 2026

Data Engineer — Michigan Campaign Finance Network

Ann Arbor, MI

- Delivered a production Flask REST API on AWS EC2 as the sole engineer on a 5-month MCFN contract, wiring ingested data and PAC rankings into the nonprofit's public-facing research workflow.
- Replaced manual Michigan campaign-finance research — portal searches capped by the Bureau of Elections, irregular Excel exports, hand normalization at 2 hours per committee — with a Requests + Pandas ETL that ingests filings directly, eliminating 800 hours of manual pulls across 400 tracked PACs.
- Architected a deterministic aggregation engine that parses those exports, normalizes irregular contribution amounts, and ranks PACs by total funding volume so researchers stop rebuilding spreadsheets to surface top spenders.

CaseStudyPrep.AI

Dec 2025 – May 2026

Software Engineer Co-op (Voice AI)

Remote

- Eliminated a 27% audio upload failure rate with fault-tolerant RxJS logic that detects expired S3 presigned URLs, regenerates them mid-flight, and negotiates MIME types for WAV files Angular silently rejected.
- Moved audio processing off the UI thread into a Web Worker with an async stream handoff, reducing main-thread blocking time to under 5ms and keeping the real-time audio visualizer at a smooth 60 FPS during active inference.

Vylet

Feb 2026 – Present | vyletdata.com

Founder — Automated lead-sourcing platform for PE/search-fund acquisition prospecting — live product generating \$1,500 MRR across three paying clients

- Launched Vylet, automating a 30-minute manual process per business into a Dockerized LangGraph pipeline generating 30 scored leads in 30 minutes — a 30x speedup — with Redis/Celery workers on a recurring cycle.
- Engineered a LangSmith eval pipeline spanning 20 adversarial business test cases across 13 archetype labels — manufacturers, SaaS tools, PE holding companies, geographic mismatches — then layered deterministic Pydantic consensus gates to lift extraction faithfulness from 50% to 90%.
- Diagnosed a name-collision defect in ownership-verification logic that was incorrectly rejecting valid acquisition targets sharing a name with an unrelated business elsewhere; the fix lifted the pipeline's lead-qualification rate from 79% to 89% with zero change to sourcing volume.

PROJECTS

SignalWeaver | *Python, FastAPI, PostgreSQL, pgvector, React, TypeScript, Docker, GitHub Actions*

Multi-signal financial research platform combining fundamentals, macro indicators, and news sentiment

- Served composite financial-research scores through async REST endpoints (FastAPI) wrapping fundamentals, MPNet sentiment, and regression logic, instrumented end-to-end at 9.1s p50 / 15.2s p99 across 90 successful runs on 90 tickers.
- Implemented semantic search over stored financial news with pgvector cosine similarity (768-d MPNet embeddings), hitting 49ms p50 / 99ms p99 over 90 queries returning top-5 or top-10 matches.
- Containerized the stack with Docker Compose (API + Postgres/pgvector + nginx frontend) and a GitHub Actions CI pipeline (frontend build, pytest, API image build on main).

TECHNICAL SKILLS

Languages Python, TypeScript, SQL
Frameworks Flask, FastAPI, React, Angular
Databases PostgreSQL, pgvector, Redis
Cloud & Infrastructure AWS (EC2, S3), Docker, Git, GitHub Actions, Celery