

Vedant Desai

(248) 704 – 4852 | verdent06@gmail.com | github.com/Verdent06 | linkedin.com/in/vedantde06

EDUCATION

University of Michigan

B.S. in Computer Science and Economics

Expected May 2028

Ann Arbor, MI

- **GPA:** 3.66 / 4.0
- **Coursework:** Data Structures & Algorithms, Intro to Statistics and Data Analysis, Microeconomics, Macroeconomics, Discrete Mathematics, Calculus III, Physics (Mechanics)

EXPERIENCE

Vylet

May 2026 – Present | vyletdata.com

Founder — Automated lead-sourcing platform for PE/search-fund; live product generating \$1,500 MRR across three clients

- Engineered a LangSmith eval pipeline spanning 20 adversarial business test cases across 13 archetype labels — manufacturers, SaaS tools, PE holding companies, geographic mismatches — then layered deterministic Pydantic consensus gates to lift extraction faithfulness from 50% to 90%.
- Launched Vylet, automating a ~30-minute manual process per business into a Dockerized LangGraph pipeline generating 30 scored leads in 30 minutes — a 30x speedup — with Redis/Celery workers on a recurring cycle.
- Architected a custom asyncpg Data Access Layer storing Gemini embeddings alongside source records; engineered injection-safe SQL timestamp validation that detects stale entries and triggers automatic re-scrapes, keeping the lead database fresh without manual intervention.

CaseStudyPrep.AI

Dec 2025 – May 2026

Software Engineer Co-op (Voice AI)

Remote

- Eliminated the silent-audio problem in a voice-AI product — most frames sent to Whisper were dead air — by running Silero VAD client-side via ONNX Runtime, filtering silence before upload and cutting cloud inference costs by 40%.
- Moved audio processing off the UI thread into a Web Worker with an async stream handoff, reducing main-thread blocking time to under 5ms and keeping the real-time audio visualizer at a smooth 60 FPS during active inference.

Lyndbrook Capital

Feb 2026 – Apr 2026

Data Engineering Consultant

Remote

- Contracted pre-LOI to build acquisition intelligence for a boutique search fund targeting water utility operators; aggregated EPA ECHO and MassGIS regulatory data into a unified PWSID entity database and delivered 800+ validated Day-1 acquisition targets within the engagement window.
- Built a Review Velocity scoring algorithm that proxied fleet expansion and operational scale from public compliance data, filtering 800 acquisition targets down to a 280-lead qualified shortlist — a 35% precision rate against the fund's revenue criteria.

PROJECTS

SignalWeaver | *Python, FastAPI, LangGraph, pgvector*

github.com/Verdent06/SignalWeaver

Multi-signal financial research platform combining fundamentals, macro indicators, and news sentiment (research assistant; not investment advice)

- Orchestrated a 5-node LangGraph pipeline (fetch → classify → embed → score → explain) with per-node timing from the same 90-run batch: fetch 7.3s p50, classify 1.1s p50 (~29ms/article over 40-article batches), embed 1.1s p50, score under 1ms p50, explain 3ms p50 — fetch dominated wall time (~80% of p50 latency).
- Implemented semantic search over stored financial news with pgvector cosine similarity (768-d MPNet embeddings), hitting 49ms p50 / 99ms p99 over 90 queries returning top-5 or top-10 matches.
- Served composite financial-research scores through async REST endpoints (FastAPI) wrapping fundamentals, MPNet sentiment, and regression logic, instrumented end-to-end at 9.1s p50 / 15.2s p99 across 90 successful runs on 90 tickers.

TECHNICAL SKILLS

Languages	Python, SQL
AI / ML	LangGraph, LangSmith, Gemini, ONNX Runtime
Frameworks	FastAPI
Databases	Redis, pgvector
Cloud & Infrastructure	Docker, Celery