

Vedant Desai

(248) 704 – 4852 | vedantde@umich.edu | github.com/Verdent06 | linkedin.com/in/vedantde06

EDUCATION

University of Michigan

B.S. in Computer Science and Economics

Expected May 2028

Ann Arbor, MI

- **GPA:** 3.66 / 4.0
- **Coursework:** Data Structures & Algorithms, Intro to Statistics and Data Analysis, Microeconomics, Macroeconomics, Discrete Mathematics, Calculus III, Physics (Mechanics)

EXPERIENCE

Michigan Data Consulting (MDC)

Jan 2026 – May 2026

Data Engineer — Michigan Campaign Finance Network

Ann Arbor, MI

- Replaced manual Michigan campaign-finance research — portal searches capped by the Bureau of Elections, irregular Excel exports, hand normalization at 2 hours per committee — with a Requests + Pandas ETL that ingests filings directly, eliminating 800 hours of manual pulls across 400 tracked PACs.
- Delivered a production Flask REST API on AWS EC2 as the sole engineer on a 5-month MCFN contract, wiring ingested data and PAC rankings into the nonprofit's public-facing research workflow.
- Scoped technical delivery directly with MCFN stakeholders as the only engineer on contract — from ingestion pipelines through REST endpoints on EC2 — with no backend team to share infrastructure, API design, or deployment ownership.

Vylet

May 2026 – Present | vyletdata.com

Founder — Automated lead-sourcing platform for PE/search-fund; live product generating \$1,500 MRR across three clients

- Launched Vylet, automating a 30-minute manual process per business into a Dockerized LangGraph pipeline generating 30 scored leads in 30 minutes — a 30x speedup — with Redis/Celery workers on a recurring cycle.
- Engineered a LangSmith eval pipeline spanning 20 adversarial business test cases across 13 archetype labels — manufacturers, SaaS tools, PE holding companies, geographic mismatches — then layered deterministic Pydantic consensus gates to lift extraction faithfulness from 50% to 90%.
- Architected a custom asyncpg Data Access Layer storing Gemini embeddings alongside source records; engineered injection-safe SQL timestamp validation that detects stale entries and triggers automatic re-scrapes, keeping the lead database fresh without manual intervention.

Lyndbrook Capital

Feb 2026 – Apr 2026

Data Engineering Consultant

Remote

- Contracted pre-LOI to build acquisition intelligence for a boutique search fund targeting water utility operators; aggregated EPA ECHO and MassGIS regulatory data into a unified PWSID entity database and delivered 800+ validated Day-1 acquisition targets within the engagement window.
- Built a Review Velocity scoring algorithm that proxied fleet expansion and operational scale from public compliance data, filtering 800 acquisition targets down to a 280-lead qualified shortlist — a 35% precision rate against the fund's revenue criteria.

PROJECTS

SignalWeaver | *Python, FastAPI, PostgreSQL, Docker, GitHub Actions* | github.com/Verdent06/SignalWeaver

Multi-signal financial research platform combining fundamentals, macro indicators, and news sentiment (research assistant; not investment advice)

- Served composite financial-research scores through async REST endpoints (FastAPI) wrapping fundamentals, MPNet sentiment, and regression logic, instrumented end-to-end at 9.1s p50 / 15.2s p99 across 90 successful runs on 90 tickers.
- Containerized the stack with Docker Compose (API + Postgres/pgvector + nginx frontend) and a GitHub Actions CI pipeline (frontend build, pytest, API image build on main).

TECHNICAL SKILLS

Languages Python, SQL
Frameworks FastAPI, Flask, LangGraph
Databases PostgreSQL, Redis, pgvector
Cloud & Infrastructure AWS (EC2, S3), Docker, Celery, Git, GitHub Actions